

Deployment Manifest - TbD V3.0

Deployment Manifest - TbD V3.0 ("Insight")

Status: Stable / Production Ready **Version:** 3.0.1 (Async Patched) **Architecture:** Cloud Native (Cloud Run + Pub/Sub + Vertex AI)

1. Directory Structure

Plaintext

tbd-v3/

```
  └── deploy_v3.ps1      # Deployment Automation (PowerShell - Updated Permissions)
  └── Dockerfile          # Container Definition (Python 3.10)
  └── requirements.txt    # Dependencies (Pinned for Stability)
  └── app/
      └── __init__.py     # Empty package marker
      └── main.py         # Router (Async Updated)
      └── schema.py       # Data Models (PAD Schema v0.2)
      └── services/
          └── __init__.py   # Empty package marker
          └── dispatcher.py # Service A: API & Pub/Sub Publisher
          └── worker.py      # Service B: Async Worker Logic
          └── pipeline.py    # V3 Orchestrator (Async + Buffer Fix)
          └── vision.py      # V3 Eyes (SSIM & Optical Flow)
          └── genai.py        # V3 Brain (Gemini 2.5 Pro)
          └── ocr.py          # Spatial OCR
          └── segment.py      # Scene Detection
```

2. Root Configuration Files

`requirements.txt`

Includes the fix for the `GenerativeModel1` import error and the NumPy binary crash.

Plaintext

```
# --- V1 Bridge Stack (Core Logic & API)
fastapi==0.104.1
uvicorn==0.24.0
pydantic==2.5.2
```

```
python-multipart
opencv-python-headless==4.8.1.78
pytesseract==0.3.10
scenedetect[opencv]==0.6.2

# --- V2 Cloud Native Stack (Infrastructure)
google-cloud-storage==2.14.0
google-cloud-pubsub==2.19.0

# --- V3 Insight Stack (New Vision & AI)
scikit-image==0.22.0
# BUMPED: Required for stable vertexai.generative_models support
google-cloud-aiplatform>=1.60.0
# PINNED: Prevent binary incompatibility with OpenCV/Scikit-Image (NumPy 2.0 crash fix)
numpy>=1.26.0,<2.0.0
```

Dockerfile

Upgraded to Python 3.10.

Dockerfile

```
FROM python:3.10-slim
```

```
# Set working directory
WORKDIR /usr/src
```

```
# Ensure Python logs are streamed directly to the terminal
ENV PYTHONUNBUFFERED=True
```

```
# 1. Install System Dependencies
```

```
# Added critical libraries for OpenCV headless and Tesseract
```

```
RUN apt-get update && \
```

```
apt-get install -y tesseract-ocr libtesseract-dev \
libgl1 libgl2.0-0 libsm6 libxext6 libxrender1 && \
rm -rf /var/lib/apt/lists/*
```

```
# 2. Copy and Install Python Dependencies
```

```
COPY requirements.txt .
```

```
RUN pip install --no-cache-dir -r requirements.txt
```

```
# 3. Copy Application Code
```

```
COPY ./app /usr/src/app
```

```
# 4. Define Environment
```

```
# PORT is injected by Cloud Run, defaulting to 8080 if missing
```

```
ENV PORT=8080
```

```
# 5. Start the Application
CMD sh -c "uvicorn app.main:app --host 0.0.0.0 --port ${PORT}"
```

deploy_v3.ps1

Updated with `objectAdmin` (to fix delete/overwrite errors) and `aiplatform.user`.

PowerShell

```
# =====
# TbD V3.0 "Insight" - Automated Deployment Script (PowerShell)
# =====
```

```
# --- CONFIGURATION ---
# REPLACE WITH YOUR ACTUAL V2 PROJECT ID
$PROJECT_ID = "tbd-v2"
$REGION = "us-central1"
$REPO_NAME = "tbd-repo"
$IMAGE_NAME = "tbd-v3-engine"
$TAG = "v3.0.1"
```

Resource Names

```
$TOPIC_NAME = "tbd-ingest-tasks"
$INPUT_BUCKET = "tbd-raw-video-$PROJECT_ID"
$OUTPUT_BUCKET = "tbd-results-$PROJECT_ID"
$DISPATCHER_SERVICE = "tbd-dispatcher"
$WORKER_SERVICE = "tbd-worker"
```

Service Accounts

```
$DISPATCHER_SA = "tbd-dispatcher-sa"
$WORKER_SA = "tbd-worker-sa"
$SUBSCRIPTION_SA = "tbd-sub-invoker"
```

```
# =====
```

```
Write-Host "--- STARTING V3 DEPLOYMENT FOR PROJECT: $PROJECT_ID ---"
```

-ForegroundColor Cyan

```
gcloud config set project $PROJECT_ID
```

1. Enable APIs (Added `aiplatform` for V3)

```
Write-Host "--- Enabling GCP APIs (including Vertex AI) ---" -ForegroundColor Green
gcloud services enable artifactregistry.googleapis.com cloudbuild.googleapis.com
run.googleapis.com pubsub.googleapis.com storage.googleapis.com aiplatform.googleapis.com
```

2. Create Storage Buckets

```
Write-Host "--- Creating GCS Buckets ---" -ForegroundColor Green
```

```

gsutil mb -l $REGION "gs://$INPUT_BUCKET/"
if ($LASTEXITCODE -ne 0) { Write-Host "Bucket might already exist, continuing..." -ForegroundColor Yellow }
gsutil mb -l $REGION "gs://$OUTPUT_BUCKET/"
if ($LASTEXITCODE -ne 0) { Write-Host "Bucket might already exist, continuing..." -ForegroundColor Yellow }

# 3. Create Pub/Sub Topic
Write-Host "--- Creating Pub/Sub Topic ---" -ForegroundColor Green
gcloud pubsub topics create $TOPIC_NAME
if ($LASTEXITCODE -ne 0) { Write-Host "Topic might already exist, continuing..." -ForegroundColor Yellow }

# 4. Build & Push Container Image
Write-Host "--- Building V3 Container Image ---" -ForegroundColor Green
gcloud artifacts repositories create $REPO_NAME --repository-format=docker
--location=$REGION --description="TbD Engine Repository"
if ($LASTEXITCODE -ne 0) { Write-Host "Repo might already exist, continuing..." -ForegroundColor Yellow }

$IMAGE_URI =
"$REGION-docker.pkg.dev/$PROJECT_ID/$REPO_NAME/$IMAGE_NAME:$TAG"

# Submit build
gcloud builds submit --tag $IMAGE_URI .

# 5. Setup Service Accounts & IAM
Write-Host "--- Configuring IAM Service Accounts ---" -ForegroundColor Green
gcloud iam service-accounts create $DISPATCHER_SA --display-name "TbD Dispatcher"
if ($LASTEXITCODE -ne 0) { Write-Host "SA might already exist..." -ForegroundColor Yellow }
gcloud iam service-accounts create $WORKER_SA --display-name "TbD Worker"
if ($LASTEXITCODE -ne 0) { Write-Host "SA might already exist..." -ForegroundColor Yellow }
gcloud iam service-accounts create $SUBSCRIPTION_SA --display-name "PubSub Invoker"
if ($LASTEXITCODE -ne 0) { Write-Host "SA might already exist..." -ForegroundColor Yellow }

# Grant Dispatcher permission to Publish
gcloud pubsub topics add-iam-policy-binding $TOPIC_NAME ` 
--member="serviceAccount:$DISPATCHER_SA@$PROJECT_ID.iam.gserviceaccount.com" ` 
--role="roles/pubsub.publisher"

# Grant Worker permission to Read Input
gsutil iam ch
"serviceAccount:$WORKER_SA@$PROJECT_ID.iam.gserviceaccount.com:objectViewer"
"gs://$INPUT_BUCKET"

```

```

# [UPDATED] Grant Worker Object Admin on Output (Allows Overwrite/Delete)
gsutil iam ch
"serviceAccount:$WORKER_SA@$PROJECT_ID.iam.gserviceaccount.com:objectAdmin"
"gs://$OUTPUT_BUCKET"

# [NEW FOR V3] Grant Worker permission to use Vertex AI (Gemini)
Write-Host "--- Granting Vertex AI User Role to Worker ---" -ForegroundColor Green
gcloud projects add-iam-policy-binding $PROJECT_ID ` 
    --member="serviceAccount:$WORKER_SA@$PROJECT_ID.iam.gserviceaccount.com" ` 
    --role="roles/aiplatform.user"

# 6. Deploy Dispatcher Service
Write-Host "--- Deploying DISPATCHER Service ---" -ForegroundColor Green
gcloud run deploy $DISPATCHER_SERVICE ` 
    --image $IMAGE_URI ` 
    --region $REGION ` 
    --service-account "$DISPATCHER_SA@$PROJECT_ID.iam.gserviceaccount.com" ` 
    --allow-unauthenticated ` 
    --memory 2Gi ` 
    --cpu 1 ` 
    --set-env-vars "SERVICE_TYPE=dispatcher,GCP_PROJECT_ID=$PROJECT_ID" ` 
    --tag v3-stable

# 7. Deploy Worker Service (V3 Insight Engine)
Write-Host "--- Deploying WORKER Service (V3 Insight Engine) ---" -ForegroundColor Green
# Increased Memory to 4Gi
gcloud run deploy $WORKER_SERVICE ` 
    --image $IMAGE_URI ` 
    --region $REGION ` 
    --service-account "$WORKER_SA@$PROJECT_ID.iam.gserviceaccount.com" ` 
    --no-allow-unauthenticated ` 
    --memory 4Gi ` 
    --cpu 2 ` 
    --timeout 3600 ` 
    --set-env-vars "SERVICE_TYPE=worker,GCP_PROJECT_ID=$PROJECT_ID" ` 
    --tag v3-stable

# 8. Create Pub/Sub Push Subscription
Write-Host "--- Linking Pub/Sub to Worker ---" -ForegroundColor Green
Write-Host "Waiting for service propagation..." -ForegroundColor Cyan
Start-Sleep -Seconds 10

```

```

$WORKER_URL = gcloud run services describe $WORKER_SERVICE --region $REGION
--format 'value(status.url)'
if (-not $WORKER_URL) {
    Write-Error "Worker Service failed to deploy or URL retrieval failed."
    exit 1
}
Write-Host "Worker URL retrieved: $WORKER_URL" -ForegroundColor Cyan

# Allow Subscription SA to invoke Worker
gcloud run services add-iam-policy-binding $WORKER_SERVICE --region $REGION `

--member="serviceAccount:$SUBSCRIPTION_SA@$PROJECT_ID.iam.gserviceaccount.com"
`

--role="roles/run.invoker"

# Create or Update Subscription
Write-Host "Configuring Subscription..." -ForegroundColor Cyan
gcloud pubsub subscriptions create tbd-worker-sub `

--topic $TOPIC_NAME `

--push-endpoint=$WORKER_URL `

--push-auth-service-account="$SUBSCRIPTION_SA@$PROJECT_ID.iam.gserviceaccount.co
m" `

--ack-deadline=600
if ($LASTEXITCODE -ne 0) {
    Write-Host "Subscription exists, updating endpoint..." -ForegroundColor Yellow
    gcloud pubsub subscriptions update tbd-worker-sub `

--push-endpoint=$WORKER_URL `

--push-auth-service-account="$SUBSCRIPTION_SA@$PROJECT_ID.iam.gserviceaccount.co
m"
}

Write-Host "====="
-ForegroundColor Cyan
Write-Host "V3 DEPLOYMENT COMPLETE" -ForegroundColor Cyan
Write-Host "Dispatcher URL: $(gcloud run services describe $DISPATCHER_SERVICE --region
$REGION --format 'value(status.url)')"
Write-Host "=====`
-ForegroundColor Cyan

```

3. App Logic Files (app/)

app/main.py

Router updated to use `await` for the worker call.

Python

```
import uvicorn
import os
from fastapi import FastAPI, HTTPException
from app.schema import TaskPayload

SERVICE_TYPE = os.environ.get("SERVICE_TYPE", "dispatcher")
app = FastAPI(title=f"TbD Engine V3 - {SERVICE_TYPE.upper()} Service")

if SERVICE_TYPE == "worker":
    try:
        from app.services.worker import WorkerService
        worker = WorkerService()
    except ImportError as e:
        print(f"CRITICAL: Failed to import WorkerService. Check dependencies. {e}")
        worker = None

    @app.post("/")
    async def pubsub_trigger(data: dict):
        if not worker:
            raise HTTPException(status_code=500, detail="Worker service failed to initialize.")
        try:
            # UPDATED: Await the worker
            await worker.process_pubsub_message(data)
            return {"status": "Processing initiated"}, 200
        except Exception as e:
            print(f"Worker processing failed: {e}")
            raise HTTPException(status_code=500, detail=f"Worker failure: {e}")

elif SERVICE_TYPE == "dispatcher":
    from app.services.dispatcher import DispatcherService
    dispatcher = DispatcherService()

    @app.post("/submit", status_code=202)
    async def submit_video_task(payload: TaskPayload):
        try:
            task_id = dispatcher.submit_task(payload)
            return {"status": "Task accepted and queued", "task_id": task_id}
        except Exception as e:
            raise HTTPException(status_code=500, detail=f"Dispatch failed: {e}")

if __name__ == "__main__":
    uvicorn.run(app, host="0.0.0.0", port=8000)
```

```
uvicorn.run(app, host="0.0.0.0", port=8080)
```

app/schema.py

V3 Schema (v0.2).

Python

```
from pydantic import BaseModel, Field
from typing import List, Optional, Dict, Any
```

--- V2 Task Payload ---

```
class TaskPayload(BaseModel):
    task_id: str = Field(..., description="UUID-V4 for the task")
    client_id: str = Field(..., description="Identifier for the client")
    gcs_uri: str = Field(..., description="gs://bucket/video.mp4")
    output_bucket: str = Field(..., description="GCS bucket for results")
    config: Dict[str, Any] = Field(default_factory=dict, description="Config params")
```

--- V3 Node Object (PAD Schema v0.2) ---

```
class ActionNode(BaseModel):
    id: str = Field(..., description="Unique node identifier")
    type: str = Field("action", description="Fixed type")
    timestamp_start: float = Field(..., description="Start time")
    timestamp_end: float = Field(..., description="End time")

    # Semantic Data (V3)
    description: Optional[str] = Field(None, description="Legacy description")
    semantic_description: Optional[str] = Field(None, description="GenAI summary")
    screen_context: Optional[str] = Field(None, description="High-level UI context")
```

Classification & Location (V3)

```
action_type: str = Field("click", description="click, type, drag, scroll")
action_class: Optional[str] = Field(None, description="navigation, input, confirmation")
```

Visual Data

```
ui_element_text: str = Field(..., description="OCR result from Active Region")
ui_region: List[int] = Field(..., description="[x, y, w, h] bounding box")
```

Confidence

```
confidence: float = Field(..., description="OCR confidence")
active_region_confidence: float = Field(0.0, description="SSIM confidence")
```

```
next_node_id: Optional[str] = Field(None, description="Next node ID")
```

--- Root Object ---

```
class Pathway(BaseModel):
```

```

pathway_id: str = Field(..., description="UUID-v4")
title: str = Field(..., description="SOP Title")
author_id: str = Field(..., description="User ID")
source_video: str = Field(..., description="Source Video URI")
created_at: str = Field(..., description="ISO-8601 Timestamp")
total_duration_sec: float = Field(..., description="Total duration")
nodes: List[ActionNode] = Field(..., description="List of nodes")
metadata: Dict[str, Any] = Field(default_factory=dict)

```

4. Service Logic Files (app/services/)

app/services/pipeline.py

The Async Orchestrator. Includes the fix for "Could not read frame at end of video".

Python

```

import os
import uuid
import time
import cv2
import asyncio
from typing import List, Tuple
from app.schema import Pathway, ActionNode
from app.services.segment import detect_action_segments
from app.services.ocr import run_ocr
from app.services.vision import detect_active_region, classify_action_optical_flow
from app.services.genai import generate_semantic_description

async def _enrich_nodes_with_ai(nodes_with_frames: List[Tuple[ActionNode,
cv2.typing.MatLike]]):
    tasks = []
    print(f"Starting async GenAI enrichment for {len(nodes_with_frames)} nodes...")
    for node, frame in nodes_with_frames:
        tasks.append(
            generate_semantic_description(
                node_id=node.id,
                frame=frame,
                ocr_text=node.ui_element_text,
                action_type=node.action_type
            )
        )
    descriptions = await asyncio.gather(*tasks)
    for i, desc in enumerate(descriptions):
        nodes_with_frames[i][0].semantic_description = desc

```

```

def _get_frame_at_time(cap: cv2.VideoCapture, timestamp: float) -> cv2.typing.MatLike:
    # FIXED: Clamp timestamp to video duration - 0.1s to prevent "read past end" errors
    max_duration = cap.get(cv2.CAP_PROP_FRAME_COUNT) / cap.get(cv2.CAP_PROP_FPS)
    safe_timestamp = min(timestamp, max_duration - 0.1)

    fps = cap.get(cv2.CAP_PROP_FPS)
    frame_no = int(safe_timestamp * fps)

    cap.set(cv2.CAP_PROP_POS_FRAMES, frame_no)
    ret, frame = cap.read()
    if not ret:
        # Fallback: Try reading the 2nd to last frame
        total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
        cap.set(cv2.CAP_PROP_POS_FRAMES, total_frames - 2)
        ret, frame = cap.read()
        if not ret:
            raise ValueError(f"Could not read frame at {timestamp}s")
    return frame

async def build_pathway(video_path: str) -> Pathway:
    print(f"Starting V3 'Insight' processing for: {os.path.basename(video_path)}")
    start_process_time = time.time()

    if not os.path.exists(video_path):
        raise FileNotFoundError(f"Video file not found: {video_path}")

    cap = cv2.VideoCapture(video_path)
    if not cap.isOpened():
        raise IOError(f"Cannot open video: {video_path}")

    fps = cap.get(cv2.CAP_PROP_FPS)
    frame_count = cap.get(cv2.CAP_PROP_FRAME_COUNT)
    total_duration_sec = frame_count / fps

    print("Running segmentation...")
    segments = detect_action_segments(video_path)

    nodes: List[ActionNode] = []
    nodes_for_enrichment: List[Tuple[ActionNode, cv2.typing.MatLike]] = []

    for i, (start_time, end_time) in enumerate(segments):
        try:
            frame_start = _get_frame_at_time(cap, start_time)
            frame_end = _get_frame_at_time(cap, end_time)

```

```

# 1. Active Region (SSIM)
active_region = detect_active_region(frame_start, frame_end)

# 2. Action Class (Optical Flow)
action_type = classify_action_optical_flow(frame_start, frame_end)

# 3. Spatial OCR
mid_time = (start_time + end_time) / 2
key_frame = _get_frame_at_time(cap, mid_time)
ocr_results = run_ocr(key_frame, active_region=active_region)

if ocr_results:
    best_result = max(ocr_results, key=lambda x: x['confidence'])
    ui_text = best_result['text']
    ui_region = best_result['ui_region']
    confidence = best_result['confidence']
else:
    ui_text = "Unlabeled Icon"
    ui_region = list(active_region) if active_region else [0,0,0,0]
    confidence = 0.0

node = ActionNode(
    id=f"node_{i + 1}",
    timestamp_start=start_time,
    timestamp_end=end_time,
    description=f"User action on '{ui_text}'",
    action_type=action_type,
    ui_element_text=ui_text,
    ui_region=ui_region,
    confidence=confidence,
    active_region_confidence=1.0 if active_region else 0.0,
    next_node_id=f"node_{i + 2}" if i + 1 < len(segments) else None
)

nodes.append(node)
nodes_for_enrichment.append((node, key_frame))

except Exception as e:
    print(f"Error processing segment {i}: {e}")
    continue

cap.release()

```

```

# UPDATED: Direct await for async GenAI
if nodes_for_enrichment:
    try:
        await _enrich_nodes_with_ai(nodes_for_enrichment)
    except Exception as e:
        print(f"GenAI Enrichment Failed: {e}")

pathway = Pathway(
    pathway_id=str(uuid.uuid4()),
    title=f"Auto-Generated Workflow: {os.path.basename(video_path)}",
    author_id="tbd-v3-engine",
    source_video=os.path.basename(video_path),
    created_at=time.strftime('%Y-%m-%dT%H:%M:%S%z'),
    total_duration_sec=total_duration_sec,
    nodes=nodes
)
print(f"V3 Processing complete in {time.time() - start_process_time:.2f}s")
return pathway

```

app/services/worker.py

Async Worker logic.

Python

```

import os
import tempfile
import base64
import json
import asyncio
from urllib.parse import urlparse
from google.cloud import storage
from app.schema import TaskPayload, Pathway
from app.services.pipeline import build_pathway

```

```
TEMP_DIR = tempfile.gettempdir()
```

```
class WorkerService:
```

```
    def __init__(self):
```

```
        try:
```

```
            self.storage_client = storage.Client()
```

```
        except Exception as e:
```

```
            print(f"WARNING: Storage client failed to initialize: {e}")
            self.storage_client = None
```

```
    def _parse_gcs_uri(self, gcs_uri: str) -> tuple:
```

```

parsed = urlparse(gcs_uri)
if parsed.scheme != 'gs':
    raise ValueError(f"Invalid GCS URI scheme: {gcs_uri}")
return parsed.netloc, parsed.path.lstrip('/')

def _download_video(self, bucket_name: str, blob_name: str, local_path: str):
    if not self.storage_client:
        print("MOCK DOWNLOAD")
        return
    bucket = self.storage_client.bucket(bucket_name)
    blob = bucket.blob(blob_name)
    print(f"Downloading gs://{bucket_name}/{blob_name} to {local_path}...")
    blob.download_to_filename(local_path)
    print("Download complete.")

def _upload_pathway(self, pathway: Pathway, output_bucket_name: str, task_id: str) -> str:
    output_blob_name = f"{task_id}/pathway.json"
    gcs_uri = f"gs://{output_bucket_name}/{output_blob_name}"
    if not self.storage_client:
        return gcs_uri
    bucket = self.storage_client.bucket(output_bucket_name)
    blob = bucket.blob(output_blob_name)
    json_data = pathway.model_dump_json(indent=2)
    blob.upload_from_string(json_data, content_type='application/json')
    print(f"Pathway uploaded successfully to: {gcs_uri}")
    return gcs_uri

# UPDATED: Async method
async def process_pubsub_message(self, pubsub_message_data: dict):
    if 'message' not in pubsub_message_data:
        raise ValueError("Invalid Pub/Sub message format.")

    message_data_bytes = base64.b64decode(pubsub_message_data['message']['data'])
    payload_dict = json.loads(message_data_bytes.decode('utf-8'))
    payload = TaskPayload.model_validate(payload_dict)

    print(f"--- Worker START: Task {payload.task_id} ---")
    input_bucket, input_blob = self._parse_gcs_uri(payload.gcs_uri)
    local_video_path = os.path.join(TEMP_DIR, os.path.basename(input_blob))

    try:
        self._download_video(input_bucket, input_blob, local_video_path)

    # UPDATED: Await the pipeline

```

```

pathway = await build_pathway(local_video_path)

output_uri = self._upload_pathway(pathway, payload.output_bucket, payload.task_id)
print(f"Worker SUCCESS: Task {payload.task_id} completed. Output URI: {output_uri}")

except Exception as e:
    print(f"Worker FAILED: {e}")
    raise e
finally:
    if os.path.exists(local_video_path):
        os.remove(local_video_path)
    print(f"--- Worker END: Task {payload.task_id} ---")

```

app/services/genai.py

Using `gemini-2.5-pro` and increased token limit.

Python

```

import os
import cv2
import vertexai
from vertexai.generative_models import GenerativeModel, Part

# Configuration
# UPDATED: Using Gemini 2.5 Pro as verified by user
MODEL_NAME = "gemini-2.5-pro"
PROJECT_ID = os.environ.get("GCP_PROJECT_ID")
LOCATION = "us-central1"

# Initialize Vertex AI
if PROJECT_ID:
    try:
        vertexai.init(project=PROJECT_ID, location=LOCATION)
    except Exception as e:
        print(f"WARNING: Vertex AI init failed: {e}")

def encode_image_to_bytes(frame):
    """Converts an OpenCV image (numpy array) to raw bytes."""
    success, encoded_image = cv2.imencode(".jpg", frame)
    if not success:
        raise ValueError("Failed to encode image")
    return encoded_image.tobytes()

async def generate_semantic_description(node_id: str, frame, ocr_text: str, action_type: str) ->
str:
    """

```

```

Generates a semantic description using Vertex AI.

"""

if not PROJECT_ID:
    return "Desc unavailable (Vertex AI not configured)."

try:
    # UPDATED: Increased max_output_tokens to 300 to prevent cut-off sentences
    model = GenerativeModel(
        MODEL_NAME,
        generation_config={"temperature": 0.0, "max_output_tokens": 300}
    )

    image_bytes = encode_image_to_bytes(frame)
    image_part = Part.from_data(data=image_bytes, mime_type="image/jpeg")

    prompt_text = (
        f"User performed '{action_type}' on '{ocr_text}'. "
        "Write a one-sentence tutorial step description. No coordinates."
    )

    response = await model.generate_content_async([image_part, prompt_text])
    return response.text.strip()

except Exception as e:
    print(f"GenAI Error on {node_id}: {e}")
    return f"User clicked '{ocr_text}' (AI description failed)"

```

app/services/vision.py

Optical Flow and SSIM Logic.

Python

```

import cv2
import numpy as np
from skimage.metrics import structural_similarity
from typing import Tuple, Optional

def detect_active_region(frame_before: cv2.typing.MatLike, frame_after: cv2.typing.MatLike) ->
Optional[Tuple[int, int, int, int]]:
    grayA = cv2.cvtColor(frame_before, cv2.COLOR_BGR2GRAY)
    grayB = cv2.cvtColor(frame_after, cv2.COLOR_BGR2GRAY)
    (score, diff) = structural_similarity(grayA, grayB, full=True)
    diff = (diff * 255).astype("uint8")
    thresh = cv2.threshold(diff, 0, 255, cv2.THRESH_BINARY_INV | cv2.THRESH_OTSU)[1]
    contours, _ = cv2.findContours(thresh.copy(), cv2.RETR_EXTERNAL,
cv2.CHAIN_APPROX_SIMPLE)

```

```

if not contours:
    return None
largest_contour = max(contours, key=cv2.contourArea)
if cv2.contourArea(largest_contour) < 100:
    return None
x, y, w, h = cv2.boundingRect(largest_contour)
padding = 10
height, width = frame_before.shape[:2]
x = max(0, x - padding)
y = max(0, y - padding)
w = min(width - x, w + (padding * 2))
h = min(height - y, h + (padding * 2))
return x, y, w, h

def classify_action_optical_flow(frame_prev: cv2.typing.MatLike, frame_curr: cv2.typing.MatLike) -> str:
    prev_gray = cv2.cvtColor(frame_prev, cv2.COLOR_BGR2GRAY)
    curr_gray = cv2.cvtColor(frame_curr, cv2.COLOR_BGR2GRAY)
    flow = cv2.calcOpticalFlowFarneback(
        prev_gray, curr_gray, None,
        pyr_scale=0.5, levels=3, winsize=15, iterations=3, poly_n=5, poly_sigma=1.2, flags=0
    )
    mean_dy = np.mean(flow[:, :, 1])
    mag, _ = cv2.cartToPolar(flow[:, :, 0], flow[:, :, 1])
    mean_mag = np.mean(mag)

    SCROLL_THRESHOLD = 2.0
    STATIONARY_THRESHOLD = 0.5

    if abs(mean_dy) > SCROLL_THRESHOLD:
        return "scroll"
    elif mean_mag < STATIONARY_THRESHOLD:
        return "click"
    else:
        return "click"

```

app/services/ocr.py

Spatial OCR Logic.

Python

```

import cv2
import pytesseract
from typing import List, Tuple, Dict, Any, Optional

def _boxes_intersect(box1: List[int], box2: List[int]) -> bool:

```

```

x1, y1, w1, h1 = box1
x2, y2, w2, h2 = box2
x_left = max(x1, x2)
y_top = max(y1, y2)
x_right = min(x1 + w1, x2 + w2)
y_bottom = min(y1 + h1, y2 + h2)
return not (x_right < x_left or y_bottom < y_top)

def run_ocr(frame: cv2.typing.MatLike, active_region: Optional[Tuple[int, int, int, int]] = None) ->
List[Dict[str, Any]]:
    ocr_data = pytesseract.image_to_data(frame, output_type=pytesseract.Output.DICT)
    results = []
    for i in range(len(ocr_data['text'])):
        text = ocr_data['text'][i].strip()
        try:
            conf_val = float(ocr_data['conf'][i])
        except:
            conf_val = 0.0
        confidence = conf_val / 100.0

        if text and confidence > 0.5:
            ui_region = [
                int(ocr_data['left'][i]), int(ocr_data['top'][i]),
                int(ocr_data['width'][i]), int(ocr_data['height'][i])
            ]
            if active_region:
                if _boxes_intersect(ui_region, active_region):
                    results.append({'text': text, 'confidence': confidence, 'ui_region': ui_region})
            else:
                results.append({'text': text, 'confidence': confidence, 'ui_region': ui_region})
    return results

```

app/services/segment.py & app/services/dispatcher.py

(Unchanged from V2 - refer to previous manifest if needed)