

V5.0 Master Deployment Manifest: Functional Activation

1. Project Directory Structure

```
tbd-v5-master/
├── deploy_v5_full.ps1      # Master Deployment Orchestrator
├── requirements.txt         # Worker/Dispatcher dependencies (V5 Final)
└── app/                      # Worker (Service B) & Dispatcher (Service A)
    ├── main.py
    ├── schema.py
    └── services/
        ├── dispatcher.py
        ├── worker.py       # V5 Execution Core (FIXED STT AWAIT)
        ├── pipeline.py     # V5 Orchestrator (FIXED 4 ARGUMENTS)
        ├── genai.py        # Gemini 2.5 Pro (FIXED 2 ARGUMENTS)
        ├── ocr.py
        ├── segment.py
        └── vision.py
    └── tbd-encoder/          # Temporal Encoder (Service C) Deployment Assets
        ├── Dockerfile
        ├── encoder_requirements.txt
        ├── lstm_model.h5
        └── encoder_app/main.py
    └── tbd-detector/          # Object Detector (Service D) Deployment Assets
        ├── Dockerfile
        ├── detector_requirements.txt
        └── detector_app/main.py
```

2. Root Configuration Files

requirements.txt (Worker/Dispatcher Dependencies)

Plaintext

```
# --- Core Framework
fastapi==0.104.1
uvicorn==0.24.0
pydantic==2.5.2
python-multipart
```

```

# --- Infrastructure
google-cloud-storage==2.14.0
google-cloud-pubsub==2.19.0
requests==2.32.5
google-auth==2.43.0

# --- V3/V4 Vision/AI Stack
opencv-python-headless==4.8.1.78
pytesseract==0.3.10
google-cloud-aiplatform>=1.60.0
numpy>=1.26.0,<2.0.0

# --- V5 Multimodality Stack (FINAL) ---
google-cloud-speech==2.25.0
pydub==0.25.1

```

Dockerfile (Worker/Dispatcher Image)

```

Dockerfile
FROM python:3.10-slim

# Set working directory
WORKDIR /usr/src
ENV PYTHONUNBUFFERED=True

# 1. Install System Dependencies
RUN apt-get update && \
    apt-get install -y tesseract-ocr libtesseract-dev libgl1 libglib2.0-0 libsm6 libxext6 libxrender1
ffmpeg && \
    rm -rf /var/lib/apt/lists/*

# 2. Copy and Install Python Dependencies
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# 3. Copy Application Code
COPY ./app /usr/src/app

# 4. Define Environment
ENV PORT=8080

CMD sh -c "uvicorn app.main:app --host 0.0.0.0 --port ${PORT}"

```

3. Worker/Encoder Application Logic

A. app/services/worker.py (V5 Execution Core)

Python

```
import os
import tempfile
import base64
import json
import asyncio
import requests
import numpy as np
from urllib.parse import urlparse
from google.cloud import storage, pubsub_v1, speech
from google.auth.transport.requests import Request
from google.oauth2 import id_token
from typing import List, Tuple, Dict, Any
from app.schema import TaskPayload, Pathway
from app.services.pipeline import build_pathway
from pydub import AudioSegment
import uuid

TEMP_DIR = tempfile.gettempdir()
# V5 CONFIGURATION CONSTANTS (Placeholder URLs MUST be set by deploy script)
OBJECT_DETECTOR_URL = "https://tbd-object-detector-XXX-uc.a.run.app"
AGENT_TOPIC_NAME = "pad-agent-tasks"
MARKETPLACE_API_URL = "https://marketplace.freefuse.com/api/v1/register"
PROJECT_ID = os.environ.get("GCP_PROJECT_ID", "tbd-v2")
AUDIO_STAGING_BUCKET = f"tbd-audio-staging-{PROJECT_ID}"

# V5 FR-01: Idempotency Check Placeholder (In-Memory)
PROCESSED_TASKS = set()

# --- V5 Helper Functions ---

def _get_auth_token(audience: str) -> str:
    """Generates an authenticated token for the target external API."""
    try:
        auth_request = Request()
        token = id_token.fetch_id_token(auth_request, audience)
        return token
    except Exception as e:
```

```

print(f"ERROR: Could not fetch auth token for {audience}: {e}")
return ""

def _extract_audio_track(video_path: str) -> str:
    """FR-06: Uses ffmpeg to extract audio and saves it locally as an MP3."""
    audio_path = os.path.join(TEMP_DIR, f"audio_{uuid.uuid4()}.mp3")
    command = f"ffmpeg -i {video_path} -vn -acodec libmp3lame -q:a 2 {audio_path}"
    os.system(command)
    return audio_path

def _upload_audio_to_gcs(local_path: str, task_id: str) -> str:
    """FR-06: Uploads the local audio file to the dedicated staging bucket for STT API."""
    storage_client = storage.Client(project=PROJECT_ID)
    audio_blob_name = f"{task_id}/audio.mp3"
    bucket = storage_client.bucket(AUDIO_STAGING_BUCKET)
    blob = bucket.blob(audio_blob_name)
    blob.upload_from_filename(local_path)
    return f"gs://{AUDIO_STAGING_BUCKET}/{audio_blob_name}"

async def _call_speech_to_text(gcs_uri: str) -> str:
    """FR-06: Implements the actual Google Cloud Speech-to-Text API call."""
    try:
        stt_client = speech.SpeechAsyncClient()
        audio = speech.RecognitionAudio(uri=gcs_uri)
        config = speech.RecognitionConfig(
            encoding=speech.RecognitionConfig.AudioEncoding.MP3,
            sample_rate_hertz=16000,
            language_code="en-US"
        )
        operation = stt_client.long_running_recognize(config=config, audio=audio)
        # FIX: The crucial 'await' is added to resolve the 'coroutine' object error
        result = await operation.result()

        transcript = " ".join([r.alternatives[0].transcript for r in result.results])
        return transcript if transcript else "No audible speech detected."
    except Exception as e:
        print(f"ERROR: STT API call failed: {e}")
        return "Audio transcription service failed."

async def _register_asset_with_marketplace(pathway_json: Dict[str, Any], trace_id: str) -> Tuple[bool, str]:
    """V5 FR-05: Makes an authenticated POST call to the external Marketplace API."""
    token = _get_auth_token(MARKETPLACE_API_URL)

```

```

if not token:
    return False, "Marketplace call failed: Missing Auth Token"

headers = {
    "Authorization": f"Bearer {token}",
    "Content-Type": "application/json",
    "X-Cloud-Trace-Context": trace_id
}

try:
    loop = asyncio.get_event_loop()
    response = await loop.run_in_executor(
        None,
        lambda: requests.post(MARKETPLACE_API_URL, headers=headers, json=payload,
    timeout=5)
    )
    response.raise_for_status()
    return True, "Marketplace registration successful."
except Exception as e:
    return False, f"Marketplace submission failed: {e}"

def _publish_to_agent_topic(gcs_uri: str, trace_id: str):
    """V5 FR-04: Publishes the final URI to the Agent Execution Topic."""
    try:
        publisher = pubsub_v1.PublisherClient()
        topic_path = publisher.topic_path(PROJECT_ID, AGENT_TOPIC_NAME)
        data = gcs_uri.encode("utf-8")
        publisher.publish(topic_path, data, trace_id=trace_id)
        print(f"✅ Published final URI to agent topic: {AGENT_TOPIC_NAME}")
    except Exception as e:
        print(f"FATAL: Agent publication failed: {e}")

# --- WorkerService Class (V5 Integration) ---

class WorkerService:
    def __init__(self):
        try:
            self.storage_client = storage.Client()
            self.agent_publisher = pubsub_v1.PublisherClient()
        except Exception as e:
            print(f"WARNING: Client initialization failed: {e}")
            self.storage_client = None
            self.agent_publisher = None

```

```

def _parse_gcs_uri(self, gcs_uri: str) -> tuple:
    parsed = urlparse(gcs_uri)
    if parsed.scheme != 'gs':
        raise ValueError(f"Invalid GCS URI scheme: {gcs_uri}")
    return parsed.netloc, parsed.path.lstrip('/')

def _download_video(self, bucket_name: str, blob_name: str, local_path: str):
    if not self.storage_client:
        print("MOCK DOWNLOAD")
        return
    bucket = self.storage_client.bucket(bucket_name)
    blob = bucket.blob(blob_name)
    print(f"Downloading gs://{bucket_name}/{blob_name} to {local_path}...")
    blob.download_to_filename(local_path)
    print("Download complete.")

def _upload_pathway(self, pathway: Pathway, output_bucket_name: str, task_id: str) -> str:
    output_blob_name = f'{task_id}/pathway.json'
    gcs_uri = f'gs://{output_bucket_name}/{output_blob_name}'

    if not self.storage_client:
        print(f"MOCK UPLOAD: Would upload result to {gcs_uri}")
        return gcs_uri

    bucket = self.storage_client.bucket(output_bucket_name)
    blob = bucket.blob(output_blob_name)
    json_data = pathway.model_dump_json(indent=2)
    blob.upload_from_string(json_data, content_type='application/json')

    print(f"Pathway uploaded successfully to: {gcs_uri}")
    return gcs_uri

async def process_pubsub_message(self, pubsub_message_data: dict):
    message_data_bytes = base64.b64decode(pubsub_message_data['message']['data'])
    payload_dict = json.loads(message_data_bytes.decode('utf-8'))
    payload = TaskPayload.model_validate(payload_dict)

    attributes = pubsub_message_data['message'].get('attributes', {})
    task_id = payload.task_id
    trace_id = attributes.get('trace_id', 'NO_TRACE_ID')

    print(f"--- Worker START: Task {task_id}, Trace {trace_id} ---")

```

```

# V5 FR-01: Idempotency Check
if task_id in PROCESSED_TASKS:
    print(f"WARNING: Task {task_id} already processed. ACK and exiting.")
    return

input_bucket, input_blob = self._parse_gcs_uri(payload.gcs_uri)
local_video_path = os.path.join(TEMP_DIR, os.path.basename(input_blob))
local_audio_path = "" # Define audio path here for finally block cleanup

# --- V5 EXECUTION CORE ---
try:
    # 1. Download Video
    self._download_video(input_bucket, input_blob, local_video_path)

    # V5 FR-06: Extract Audio Track -> Upload -> Transcribe
    local_audio_path = _extract_audio_track(local_video_path)
    gcs_audio_uri = _upload_audio_to_gcs(local_audio_path, task_id)
    audio_transcript = await _call_speech_to_text(gcs_audio_uri)

    # 2. Execute V4 Pipeline (Pass Transcript, Object Detector URL)
    from app.services.pipeline import build_pathway

    # FIX: The Worker now passes 4 arguments to the pipeline function
    pathway = await build_pathway(
        local_video_path,
        payload.gcs_uri,
        audio_transcript,
        OBJECT_DETECTOR_URL
    )

    # 3. Finalization and Execution Hooks
    output_uri = self._upload_pathway(pathway, payload.output_bucket, payload.task_id)

    # V5 FR-05: Call Marketplace API
    market_success, market_message = await
    _register_asset_with_marketplace(pathway.model_dump(), trace_id)

    # V5 FR-04: Publish to Agent Topic
    _publish_to_agent_topic(output_uri, trace_id)

    # 4. Final Audit and Idempotency Lock
    PROCESSED_TASKS.add(task_id)
    print(f"Worker SUCCESS: Task {task_id} completed. Output URI: {output_uri}")
    print(f"Marketplace Status: {market_message}")

```

```

except Exception as e:
    print(f"Worker FAILED: {e}")
    raise e
finally:
    if os.path.exists(local_video_path):
        os.remove(local_video_path)
        # Clean up extracted audio file
    if os.path.exists(local_audio_path):
        os.remove(local_audio_path)
print(f"--- Worker END: Task {task_id}, Trace {trace_id} ---")

```

C. app/services/pipeline.py (V5 Orchestrator)

Python

```

import os
import uuid
import time
import cv2
import asyncio
import json
import base64
import requests
from typing import List, Tuple
from app.schema import Pathway, ActionNode
from app.services.genai import analyze_video_native
from app.services.ocr import run_ocr
from google.oauth2 import id_token
from google.auth.transport.requests import Request

# V5 NEW CONFIGURATION CONSTANTS (Passed from worker)
OBJECT_DETECTOR_TIMEOUT = 5.0

def _get_auth_token(audience: str) -> str:
    """Generates an authenticated token for the target Cloud Run service."""
    try:
        auth_request = Request()
        token = id_token.fetch_id_token(auth_request, audience)
        return token
    except Exception as e:
        print(f"ERROR: Could not fetch auth token for Encoder: {e}")
        return ""

def _get_frame_at_time(cap: cv2.VideoCapture, timestamp: float) -> cv2.typing.MatLike:

```

```

"""Extracts a frame at a specific timestamp for coordinate refinement."""
max_duration = cap.get(cv2.CAP_PROP_FRAME_COUNT) / cap.get(cv2.CAP_PROP_FPS)
safe_timestamp = min(timestamp, max_duration - 0.1)

fps = cap.get(cv2.CAP_PROP_FPS)
frame_no = int(safe_timestamp * fps)

cap.set(cv2.CAP_PROP_POS_FRAMES, frame_no)
ret, frame = cap.read()
if not ret:
    total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
    cap.set(cv2.CAP_PROP_POS_FRAMES, total_frames - 2)
    ret, frame = cap.read()
if not ret:
    return None
return frame

async def _call_object_detector(frame: cv2.typing.MatLike, target_text: str, detector_url: str) ->
    Tuple[List[int], float]:
    """
    FR-07: Calls Service D securely for pixel-accurate coordinate prediction.
    """
    auth_token = _get_auth_token(detector_url)

    if frame is None or not auth_token:
        print("WARNING: Object Detector call failed. Missing auth token or frame.")
        return [0, 0, 0, 0], 0.0

    success, buffer = cv2.imencode('.jpg', frame)
    frame_base64 = base64.b64encode(buffer).decode('utf-8')

    headers = {
        "Authorization": f"Bearer {auth_token}",
        "Content-Type": "application/json"
    }
    payload = {"frame_base64": frame_base64, "target_text": target_text}

    try:
        loop = asyncio.get_event_loop()
        response = await loop.run_in_executor(
            None,
            lambda: requests.post(f"{detector_url}/detect_coordinates", headers=headers,
            json=payload, timeout=OBJECT_DETECTOR_TIMEOUT)
        )
    
```

```

response.raise_for_status()

result = response.json()
return result.get('ui_region', [0, 0, 0, 0]), result.get('confidence', 0.0)

except Exception as e:
    print(f"WARNING: Object Detector call failed: {e}")
    return [0, 0, 0, 0], 0.0

# FINAL V5 PIPELINE SIGNATURE: Correctly accepts 4 arguments
async def build_pathway(local_video_path: str, gcs_video_uri: str, audio_transcript: str,
object_detector_url: str) -> Pathway:
    print(f"Starting V4 'Native Insight' for: {os.path.basename(local_video_path)}")
    start_time = time.time()

    print("Phase 1: Semantic Analysis (Gemini Native Video + Audio Fusion)...")

    ai_steps = await analyze_video_native(gcs_video_uri, audio_transcript)
    print(f"Gemini identified {len(ai_steps)} steps.")

    print("Phase 2: Coordinate Refinement (Service D Call)...")

    cap = cv2.VideoCapture(local_video_path)

    fps = cap.get(cv2.CAP_PROP_FPS)
    total_frames = cap.get(cv2.CAP_PROP_FRAME_COUNT)
    total_duration_sec = total_frames / fps if fps else 0

    final_nodes = []

    for i, step in enumerate(ai_steps):
        timestamp = float(step.get('timestamp', 0.0))
        target_text = step.get('target_text', "Unlabeled")

        frame = _get_frame_at_time(cap, timestamp)

        ui_region, confidence = await _call_object_detector(frame, target_text, object_detector_url)

        # Construct the Node
        node = ActionNode(
            id=f"node_{i + 1}",
            timestamp_start=timestamp,
            timestamp_end=timestamp + 1.0,
            description=step.get('description', 'No description'),
            semantic_description=step.get('description', 'No description'),

```

```
        ui_element_text=target_text,
        ui_region=ui_region,
        confidence=confidence,
        active_region_confidence=confidence,
        action_type=step.get('action_type', 'click'),
        next_node_id=f"node_{i + 2}" if i + 1 < len(ai_steps) else None
    )
    final_nodes.append(node)

cap.release()

pathway = Pathway(
    pathway_id=str(uuid.uuid4()),
    title=f"Native Insight: {os.path.basename(local_video_path)}",
    author_id="tbd-v4-engine",
    source_video=os.path.basename(local_video_path),
    created_at=time.strftime('%Y-%m-%dT%H:%M:%S%z'),
    total_duration_sec=total_duration_sec,
    nodes=final_nodes
)
print(f"V4 Processing complete in {time.time() - start_time:.2f}s")
return pathway
```