

# **Product Vision Document: The Teach by Doing (TbD) Ingestion Engine**

**Product:** Teach by Doing (TbD) / PAD Ingestion Core

**Version:** 1.0 (Nov 2025 Baseline)

**Author:** Greg Burns, Product Owner

**Purpose:** To define the single, critical component required to solve the Data Gap and launch the full Pathways as Data (PAD) ecosystem.

## **1. Introduction: The Core Thesis**

The TbD engine is the **on-ramp for all knowledge** into the Pathways as Data system. It transforms an expert's recorded action (video/screen capture) into a structured, machine-readable asset called a "Pathway."

### **The Dual-Use Architecture**

The core value of PAD is that the pathway asset is immediately usable by two entities:

- For Humans:** It generates **zero-cost, human-readable documentation** (like an instant recipe from a recorded video).
- For AI:** It is an **Executable Instruction Set** that a robot, co-pilot, or AI agent can follow perfectly.

---

## **2. The Problem We Are Solving (The Gaps)**

TbD directly targets the three fundamental gaps blocking the PAD vision:

<b>Gap</b>	<b>Problem Statement</b>	<b>Example of Failure</b>
<b>Data Gap</b>	Expert knowledge is siloed, making it impossible to acquire labeled training data at scale.	DED manufacturing scrap rates (10-30%) due to inconsistent human expertise.

<b>Tech Gap</b>	Automation requires specialized (and expensive) AI/robotics engineers to manually program workflows.	Surgeons and technicians cannot train AI; they must rely on external programmers (Prompt Literacy Gap).
<b>Strategy Gap</b>	Without a structured data format, the entire vision remains theoretical and cannot be scaled or monetized.	We cannot build a "Pathway Marketplace" without a consistent, verifiable product to sell.

### 3. Product Architecture: From Video to Structure (V1-V4)

The TbD engine is built on a scalable, cloud-native architecture that moves from simple observation to intelligent, temporal understanding.

#### 3.1 The Scalable Foundation (V2.0 Baseline)

- **Decoupled Architecture:** The system uses a **Fan-Out Architecture** on GCP (Cloud Run, Pub/Sub, GCS) to safely process long-running videos without timeouts. This is now fully deployed and stable.
- **The Bridge Stack:** V1/V2 utilize a standard Computer Vision (CV) stack ([OpenCV](#), [Tesseract](#)) as a temporary "bridge" to generate initial data, bypassing the stalled external VCU dependency.

#### 3.2 The Intelligence Core (V3.0/V4.0 Focus)

The core value is added by advanced processing layers:

Layer	Function	Technical Rationale
<b>V3.0: Semantic Enrichment</b>	Replaces simple text with human-readable descriptions.	Uses <b>SSIM Differencing</b> for precise "Active Region" detection and <b>Gemini 2.5 Pro</b> to generate <b>zero-cost human documentation</b> .

<b>V4.0: Temporal Grounding</b>	<b>Solves the VCU Dependency.</b> Uses a custom <b>LSTM/RNN Temporal Encoder Microservice</b> to enforce sequential logic and <i>remember</i> the preceding steps.	This is the "Temporal Grounding Bridge" that makes the asset executable and necessary for complex actions like "baking a cake".
---------------------------------	--	---

## 4. Business & Monetization Hooks

The TbD engine is the foundation for all revenue in the PAD hybrid model:

### A. The ARM-Style IP Strategy

- Every Pathway generated by TbD is a piece of licensable intellectual property.
- The final JSON structure includes mandatory **Marketplace Metadata** (e.g., `license_tier`, `compliance_tag`) to ensure it's ready for commercialization.

### B. The Compliance & Data Moat

- TbD is built with "**Compliance-by-Design.**" This is critical for the Medical vertical.
- The system creates the **fully auditable digital log** required by standards like AS9100 (Manufacturing) and HIPAA (Medical), turning a regulatory burden into a valuable data asset.
- The ingestion pipeline is the single source of data for the "compendium" or **Data Moat**, which is the long-term competitive advantage.

## 5. Strategic Roadmap (2026 Baseline)

The original 2025 timeline is now an internal historical record<sup>1</sup>. This roadmap reflects a realistic, accelerated timeline starting with the core build and moving toward full agent integration.

Phase	Timeframe	Key Milestone	Strategic Value
<b>Phase 1: Ingestion Gateway</b>	<b>Q1 2026</b>	<b>TbD V4.0 Launch.</b> Full GCP decoupling, Semantic Enrichment, and <b>Temporal Grounding Bridge</b> deployment.	<b>Control the Timeline:</b> Unblocks the PAD roadmap from the external VCU dependency.

<b>Phase 2: Agent Execution Core</b>	<b>Q2 2026</b>	<b>VCU Final Integration (External or Internal V5.0).</b> Formal launch of the <b>Risk Agent</b> and <b>Simulation Agent</b> .	<b>Launch the Brain:</b> Activates the predictive and simulation capabilities required for the Manufacturing MVP.
<b>Phase 3: Cross-Domain Readiness</b>	<b>Q3 2026</b>	<b>Final Sensor Bridge.</b> Integration with the <b>IoT-Action Bridge</b> (MQTT/WebSockets) to accept <b>DED machine telemetry</b> data.	<b>Physical World Integration:</b> Moves beyond software to monitor and control industrial/exoskeleton processes.
<b>Phase 4: Marketplace Launch</b>	<b>Q4 2026</b>	<b>Pathway Marketplace API Finalization.</b> Finalize licensing, delivery, and reporting hooks (Ankush's domain).	<b>Monetization:</b> Activates the full royalty and revenue-share model across all verticals.