

This is a decisive move, Burnsie. You're shifting from waiting for direction to **defining the product** and providing the team with tangible assets they can use immediately.

Here is the Technical Design Document (TDD) for the V1 Teach-by-Doing (TbD) MVP. It is structured to be handed directly to a developer (or used by you) to build the system. It includes the data schema, the architecture, and the implementation plan.

Technical Design Document: Teach by Doing (TbD) MVP - V1.0

Status: Ready for Build

Author: Greg Burns

Date: November 21, 2025

Objective: To build a local, standalone prototype that ingests a screen recording of a user workflow and outputs a structured Pathway.json file, serving as the foundational "Data Ingestion Engine" for the FreeFuse PAD initiative.

1. System Overview

The TbD MVP is a **Screen-to-Schema** pipeline. It does not require real-time streaming or complex cloud infrastructure for V1. It focuses on the core problem: **Observation \rightarrow Structure**.

1.1 The "Bridge" Strategy

To bypass the delay in the "TimeSuite" VCU project, this V1 uses a "Bridge Stack" of readily available, open-source CV tools to achieve immediate functionality.

- **Input:** A raw .mp4 video file of a user performing a task on their screen (e.g., navigating a software interface).
- **Core Logic:**
 1. **Segmentation:** Break video into "action events" based on visual changes.
 2. **Extraction:** Use OCR (Tesseract) to read UI text and Template Matching (OpenCV) to identify clicks/icons.
 3. **Structuring:** Map these events to the **Pathway** data schema.
- **Output:** A JSON file that machines (agents) and humans can read.

2. Data Schema Definition (The "Contract")

This schema is the most critical artifact. It defines the "common language" between the human expert and the AI agent.

2.1 The Root Object

JSON

```
{  
  "pathway_id": "UUID-v4",  
  "metadata": {  
    "title": "Standard Operating Procedure: [Task Name]",  
    "author_id": "user_123",  
    "source_video": "s3://bucket/uploads/raw_video.mp4",  
    "created_at": "ISO-8601 Timestamp",  
    "total_duration_sec": 145.5  
  },  
  "nodes": []  
}
```

2.2 The Node Object (The Atomic Unit)

Every detected event is a "Node." In V1, we focus on **ActionNodes**.

JSON

```
{  
  "id": "node_1",  
  "type": "action",  
  "timestamp_start": 12.5,  
  "timestamp_end": 14.2,  
  "description": "User clicked 'File' menu",  
  "data": {  
    "action_type": "click", // enum: click, type, drag, scroll  
    "ui_element_text": "File", // OCR result  
    "ui_region": [10, 20, 100, 50], // [x, y, w, h] bounding box  
    "confidence": 0.92  
  },  
  "next_node_id": "node_2"  
}
```

3. Architecture & Tech Stack

This architecture is designed for local development now, with a clear path to GCP deployment (Cloud Run/Vertex AI) later.

3.1 The Stack

- **Language:** Python 3.9+

- **Video Processing:** OpenCV (`cv2`) for frame extraction and difference hashing.
- **OCR Engine:** Tesseract 5 (via `pytesseract`) for reading screen text.
- **Scene Detection:** PySceneDetect for splitting the video into logical "steps."
- **API Framework:** FastAPI (to serve the processing logic).
- **Container:** Docker (for portability).

3.2 The Pipeline Flow (The Algorithm)

Step 1: Ingest & Pre-process

- Upload `.mp4`.
- Convert to 1fps (frames per second) to reduce compute load for V1.

Step 2: Event Segmentation (The "When")

- Use PySceneDetect with `ContentDetector` algorithm.
- *Logic:* If pixel change > Threshold (30%), mark as a "New Scene" (new step).
- *Output:* A list of (`start_time`, `end_time`) tuples.

Step 3: Visual Extraction (The "What")

- For each scene, extract the **middle frame** (the stable state).
- Run **OCR** on the frame to get all text.
- *Heuristic:* If a specific small region changed significantly compared to the previous scene, assume that is the "Active Region" (e.g., a button press).

Step 4: Node Generation

- Create a `Node` object for each scene.
- Populate `ui_element_text` with the OCR text found in the "Active Region."
- Link nodes sequentially (`node_1.next -> node_2`).

4. Implementation Plan (Sprint 0)

This is the build order for your weekend sprint.

Phase 1: The Skeleton (Friday Night)

- [] Set up Python environment (`requirements.txt`: `fastapi`, `opencv-python`, `pytesseract`, `scenedetect`).
- [] Write `schema.py`: Pydantic models defining the JSON structure above.
- [] Write `ingest.py`: A script that accepts a file path and verifies it's a valid video.

Phase 2: The "Eyes" (Saturday)

- [] Implement `segmentation.py`: Use PySceneDetect to return a list of timestamps.

- [] Implement `ocr.py`: A function that takes a frame (image array) and returns text + bounding boxes.
- [] **Test:** Feed it a video of you clicking a "Save" button. Verify it outputs the text "Save".

Phase 3: The "Brain" & Output (Sunday)

- [] Implement `pipeline.py`: The orchestrator that loops through segments, calls OCR, and builds the JSON.
 - [] Write `main.py`: The FastAPI entry point (`POST /process-video`).
 - [] **Demo Artifact:** Run the script on your Surgical PoC usage video. Generate the `output.json`.
-

5. Deployment Strategy (GCP Ready)

Since you mentioned deploying to GCP:

1. **Dockerize:** Create a `Dockerfile` that installs system dependencies (`libgl1`, `tesseract-ocr`) and your Python requirements.
2. **Cloud Run:** Deploy the container to Google Cloud Run (Generation 2 for longer timeouts).
3. **Storage:** Use GCS buckets for video upload/storage.